

LAMP: A Selective-ACK Transport Protocol for Reliable Bulk Data Transmission over LoRa Networks

Alessandro Monticelli

Abstract—This paper presents Lightweight Addressed Messaging Protocol (LAMP), a lightweight transport architecture that enables the continuous transfer of large data blocks over point-to-point Long-Range (LoRa[®]) networks. To overcome the prohibitive Time-on-Air overhead and duty-cycle saturation typical of existing Stop-and-Wait solutions, the proposed architecture segments payloads into up to 255 chunks per transmission window. Each chunk is protected by a 9-byte Logical Header with compact 8-bit node addressing (supporting downstream IPv6 link-local prefix mapping) and a 2-byte Cyclic Redundancy Check (CRC)-16 checksum suffix. The protocol provides four distinct operational modes combining connection state (Connectionless / Stateful) and reliability (Best-effort / SACK-ARQ feedback). By requiring receiver feedback only at sequence boundaries rather than per-packet, the Selective Acknowledgment (SACK) mechanism significantly reduces acknowledgment overhead without incurring the constant over-the-air encoding overhead of pure Forward Error Correction (FEC).

Index Terms—Long Range (LoRa), SACK, packet fragmentation, transport protocol, IoT, embedded systems, 3-Way Handshake, point-to-point.

I. INTRODUCTION

LOW-POWER, low-cost, and long-range wireless communication technologies are becoming increasingly important in the context of the Internet of Things (IoT). Applications such as environmental monitoring, precision agriculture, and autonomous Unmanned Aerial Vehicles (UAVs) require reliable data transmission over extended distances under strict power and hardware constraints [1].

Several communication standards, including ZigBee, Sigfox, and Long Range (LoRa) have been adopted to address these needs, and LoRa, in particular, has emerged as the leading choice for low-power and long-range deployments thanks to its robust modulation technique and open network architecture [2]. However, conventional solutions based on LoRa (particularly those following the LoRaWAN[®] specifications) are optimized for small, non-frequent payloads [3].

This limitation poses acute restrictions on data-intensive IoT applications which demand the continuous transfer of large data blocks [4], [5], such as Firmware Update Over The Air (FUOTA), Simultaneous Localization and Mapping (SLAM) map transfers in autonomous robotics, or bulk sensor log collection. Standard Long Range Wide Area Network (LoRaWAN) links require substantial transmission overhead, restrictive duty-cycle limitations [6], and lack optimized mechanisms for handling large, multi-packet payloads in point-to-point configurations.

This paper presents the design and implementation of LAMP, a lightweight transport protocol for point-to-point and local multi-node LoRa networks. LAMP supports payloads up to 62.2 KB on 255 byte First-In, First-Out (FIFO) transceivers (such as the Semtech[®] SX1262), operating independently of higher-layer middleware such as LoRaWAN.

The main contributions of this work are:

- **Bitmap-Based SACK Scheme:** Unlike existing Low-Power Wide-Area Network (LPWAN) fragmentation solutions that rely on Stop-and-Wait (S&W) Automatic Repeat Request (ARQ), which requires one acknowledgment per packet and incurs substantial latency under duty-cycle constraints, the proposed protocol uses a compact binary bitmap carried in a single feedback frame. The receiver reports the delivery status of up to 255 chunks per round-trip, reducing retransmission overhead and over-the-air footprint.
- **Lightweight Packet Segmentation and Abstraction:** The protocol introduces a 9-byte logical header (with 8-bit local node addressing) and a custom CRC integrity check on partial payloads, enabling out-of-order chunk reassembly on resource-constrained microcontrollers.
- **Modular C++17 Implementation and Test-Driven Development (TDD) Validation:** The core state machine is completely decoupled from the physical radio driver through a Hardware Abstraction Layer (HAL). This design enables both native host-based unit testing (using the Unity framework) and physical deployment on Semtech SX126x transceivers.

The remainder of this paper is organized as follows. Section II surveys related work and defines the design space. Section III details the system architecture, packet structure, and formal protocol algorithms. Section IV outlines preliminary prototyping findings. Section V presents the hardware-level C++17 implementation. Section VI evaluates performance through extensive ns-3 simulations and planned field validation. Section VII concludes the paper and presents the architectural roadmap.

II. STATE OF THE ART AND DESIGN SPACE

Reliable packet fragmentation and transport over LPWANs has been the subject of extensive investigation, driven by the expanding demand for data-intensive IoT workloads [4], [5]. Existing approaches broadly fall into three categories: standardized LPWAN frameworks, conventional Stop-and-Wait ARQ, and specialized windowed batch transports.

A. Standardized LPWAN Frameworks: SCHC and LoRaWAN TS004

At the standardization level, the IETF defined Static Context Header Compression (SCHC) [7], [8], introducing a generic architecture for header compression and fragmentation over LPWANs. RFC 9011 defines SCHC profiles over LoRaWAN supporting No-ACK, ACK-Always, and ACK-on-Error modes. While SCHC achieves high compression for predictable IPv6/UDP flows (down to 1–2 bytes), it relies on pre-provisioned static rule contexts and gateway-centric infrastructure, making it ill-suited for autonomous, ad-hoc raw-LoRa peer-to-peer or multi-node links. Similarly, the LoRa Alliance standardized Fragmented Data Block Transport (TS004! (TS004!)) [9] to support FUOTA via open-loop FEC redundancy codes. While open-loop FEC delivers multicast firmware images without requiring a reverse feedback channel, it introduces continuous over-the-air encoding overhead regardless of physical channel conditions and imposes substantial computational decoding demands on resource-constrained microcontrollers.

B. Point-to-Point Transports and Stop-and-Wait Limitations

In infrastructure-free point-to-point and local mesh configurations, early implementations relied on Stop-and-Wait (S&W) ARQ. However, empirical link evaluations by Cattani et al. [10] and network scalability analyses by Bor et al. [11] established that per-fragment acknowledgments incur prohibitive Time-on-Air latencies and severe duty-cycle saturation when applied to multi-packet payloads. Under typical LPWAN channels, the round-trip overhead of waiting for an individual ACK after every single packet dramatically reduces effective throughput.

C. Windowed Selective Transports: MPLR and Technical Gaps

To mitigate Stop-and-Wait overhead, Chen et al. [12] proposed Multi-Packet LoRa (MPLR), introducing windowed batch transmissions acknowledged by a bit-vector acknowledgment packet (BVACK). While pioneering the use of bitmap feedback over LoRa, MPLR was engineered as an application-specific pipeline tailored strictly for uploading agricultural camera images from an end-device to a central gateway. Consequently, MPLR exhibits key architectural limitations when generalized as a generic transport protocol:

- 1) *Substantial Header Overhead*: MPLR enforces a 16-byte fixed header per fragment (allocating 8 bytes for 64-bit EUI source/destination addresses), which consumes a substantial portion of the 255-byte maximum LoRa physical payload.
- 2) *Rigid Single-Mode Operation*: MPLR operates solely under a stateful batch model with a fixed window size (40 packets), lacking support for unacknowledged streaming, connectionless datagrams, or dynamic window scaling.
- 3) *Static Setup and Heavy Teardown*: MPLR lacks run-time buffer negotiation and mandates an explicit 4-way

handshake teardown sequence (FIN/ACK), generating unnecessary control frame traffic.

D. Architectural Contributions of LAMP

LAMP addresses these gaps by establishing an autonomous, lightweight, and configurable transport layer operating directly over raw LoRa PHY. As summarized in Table I, LAMP advances the state of the art through four core design principles:

- **Compact Header Architecture**: LAMP utilizes a 9-byte Logical Header (+2-byte CRC-16 suffix, totaling 11 bytes over the air), reducing per-fragment control overhead by 31.2% compared to MPLR (16 bytes) while providing 8-bit local unicast/broadcast addressing.
- **Four-Mode Operational Flexibility**: Unlike single-mode batch protocols, LAMP spans a complete operational matrix: *Unreliable Connectionless* (Mode 1), *Unreliable Connected* (Mode 2), *Reliable Connectionless* (Mode 3), and *Reliable Connected* (Mode 4).
- **Dynamic 3WHS Negotiation & Safe Memory Bounds**: Incorporates run-time parameter negotiation (SynMetadata) for buffer safety on low-power microcontrollers, replacing explicit 4-way teardown overhead with automatic idle session pruning.
- **Transport-Layer Purity**: Completely decoupled from physical radio drivers via a clean HAL, enabling native host testing and embedded deployment.

III. PROPOSED SYSTEM ARCHITECTURE

LAMP is a transport-layer protocol running directly over the LoRa physical layer. It segments arbitrarily large payloads (within the described constraints) at the transmitter and reconstructs them at the receiver while minimizing Logical Header and Over-the-Air (OTA) overhead.

A. Protocol Specification and Design Iterations

We describe the design of LAMP by contrasting its initial stateless baseline (LAMPv1) with the current addressed architecture (LAMPv2).

1) *Stateless Header Specification*: The initial protocol version, LAMPv1, was designed for maximum header compression during periodic sensor telemetry transfers over point-to-point links. It utilized a compact 7-byte Logical Header containing only sequence tracking and payload bounds (messageId [2 B], totalChunks [1 B], chunkIndex [1 B], payloadSize [1 B], flags [1 B], and protocolVersion [1 B]).

While LAMPv1 achieved high spectral efficiency for isolated datagram bursts, deployment on embedded hardware revealed three key architectural limitations:

- *Speculative Buffer Allocation*: Receivers allocated re-assembly memory speculatively upon receiving the first fragment of an unannounced stream. When receiving large payloads (e.g., image tiles or firmware blocks), memory-constrained microcontrollers risked heap exhaustion or buffer overflows.

TABLE I
ARCHITECTURAL COMPARISON AND TRADE-OFF MATRIX OF LoRa TRANSPORT SOLUTIONS

Architectural Dimension	Stop-and-Wait ARQ	SCHC [7]	LoRaWAN Fragmented [9]	MPLR [12]	Proposed (LAMP)
Stack Dependency	Raw LoRa PHY	Generic Adaptation (often LoRaWAN/IP)	LoRaWAN App Layer	Raw LoRa PHY	Raw LoRa PHY (Standalone)
ACK Mechanism	Per-Packet ACK	ACK-on-Error / ACK-Always	None (Open-Loop Redundancy)	Bit-Vector SACK	Bit-Vector SACK
Header Overhead	2–4 Bytes	1–2 Bytes (Compressed RuleID/Residue)	2 Bytes	16 Bytes	9 Bytes (11B total w/ CRC)
Max Frame Payload	240+ Bytes	Variable	Variable	239 Bytes	244 Bytes (255B PHY limit)
Max Session Payload	Unbounded (Sequential)	Variable (IPv6 Packet)	Unlimited (Block Session)	Multi-MB Image (65k Packets)	62.2 KB (255 Chunks / Session)
Operational Modes	1 (Reliable S&W)	3 (No-ACK, ACK-Always, ACK-on-Error)	1 (Open-Loop FEC)	1 (Reliable Batch)	4 Configurable Modes
Node Addressing	None / Application	IPv6 128-bit (Compressed)	DevAddr (32-bit)	EUI (64-bit)	8-bit Unicast + Broadcast
Multi-Node Suitability	Single P2P Link	Star (Gateway-Centric)	Star Broadcast (Downlink)	P2P Uplink (Node-GW)	P2P & Multi-Node Ready
Parameter Negotiation	None (Static)	Static (Pre-shared Rules)	Static / Out-of-Band	Static Batch / 4WHS	Dynamic 3WHS Metadata
Key Strengths	Minimal header size, trivial implementation	Header compression (1–2 B) for IP traffic	No feedback channel required (pure FEC)	High-volume image uplink batching	4 flexible modes, dynamic 3WHS buffer safety, raw LoRa autonomy
Key Limitations & Trade-offs	Prohibitive airtime latency, duty-cycle saturation	Requires shared rule context & IP infrastructure	High FEC compute cost, constant airtime overhead	16B header overhead, single mode, 4WHS teardown	9B header (11B w/ CRC), 62.2 KB session limit, RAM/timer state

- *Session Interference*: Lacking explicit source and destination identifiers, multiple unaddressed transmitters sharing the same RF channel risked interleaving fragments with identical sequence numbers, leading to corrupted reassembly buffers.
- *Static Transport Limits*: Fixed chunk sizes and static timeout thresholds prevented dynamic adaptation to link quality degradation or variable receiver memory limits.

2) *Addressed Logical Header Specification*: To resolve the limitations of LAMPv1, LAMPv2 introduces an addressed, stateful transport architecture. LAMPv2 expands the Logical Header to 9 bytes by embedding explicit 8-bit Source (`srcAddr`) and Destination (`dstAddr`) Node Identifiers, while incorporating an integrated 3WHS mechanism.

Although adding 2 bytes to the header increases over-the-air frame length by less than 0.8% relative to the maximum 255-byte frame (9-byte header + 244-byte payload + 2-byte CRC-16), this trade-off provides three main architectural benefits:

- 1) *Buffer Memory Safety*: The 3WHS exchange allows nodes to negotiate reassembly buffer limits (`SynMetadata`) prior to payload transmission, preventing allocation failures on memory-constrained receivers.
- 2) *Channel Session Isolation*: Explicit node addressing prevents session crosstalk and frame interleaving in multi-node wireless topologies.
- 3) *IPv6 Integration Readiness*: The 8-bit Node ID introduces a lightweight local addressing namespace that can be mapped directly to an IPv6 link-local prefix at an edge gateway, enabling transparent IP-layer integration without transmitting heavy 128-bit IPv6 headers over

raw LoRa radio links.

B. Operational Modes and Sequence Management

To accommodate diverse IoT workload requirements, LAMP defines four distinct operational modes across two primary dimensions: *Reliability* (Unreliable Best-Effort vs. SACK-ARQ) and *Connection State* (Connectionless Handshake-Free vs. Connection-Oriented Stateful).

1) *Implicit Frame Boundary Derivation*: To maintain a compact 9-byte header while supporting robust out-of-order packet reassembly, LAMP avoids dedicated boundary control bits. Instead, Start-of-Message (SOM) and End-of-Message (EOM) boundaries are implicitly derived at the receiver directly from the 0-based fragment index (`chunkIndex`) and the total fragment count (`totalChunks`):

$$\text{SOM} \iff \text{chunkIndex} == 0 \quad (1)$$

$$\text{EOM} \iff \text{chunkIndex} == \text{totalChunks} - 1 \quad (2)$$

Explicitly carrying `totalChunks` in every fragment header allows the receiver to determine the exact message bounds and pre-allocate the required reassembly buffer and SACK bitmap upon receiving *any* fragment, even if fragments arrive out of order or if the initial chunk is lost.

2) *Node Addressing Scheme*: To support multi-node communication without incurring the prohibitive airtime overhead of full 128-bit IPv6 headers on LoRa PHY payloads, LAMP introduces a compact 8-bit Node Identifier space:

- **0x00 (ADDRESS_UNASSIGNED)**: Reserved for legacy or unaddressed P2P datagrams (promiscuous mode).
- **0x01–0xFE**: 254 unique individual node addresses (Unicast).

TABLE II
LAMP OPERATIONAL MODES ARCHITECTURE

Operational Mode	Reliability & Connection State	Primary Target Use Case
Mode 1: Best-Effort Datagram	Unreliable, Connectionless (Stateless)	High-frequency telemetry streams, periodic sensor data bursts.
Mode 2: Best-Effort Stream	Unreliable, Connection-Oriented (Stateful)	Session-tracked real-time sensor streams without retransmission delays.
Mode 3: SACK-Reliable Datagram	Reliable (SACK), Connectionless (Handshake-Free)	Isolated high-priority sensor bursts requiring delivery confirmation without prior session setup.
Mode 4: SACK-Reliable Stream	Reliable (SACK), Connection-Oriented (3WHS)	Critical large payload transfers (FUOTA binaries, camera images, SLAM maps).

- **0xFF (ADDRESS_BROADCAST)**: Broadcast address processed by all active receiving nodes on the RF channel. As detailed in Section III-A, this 8-bit Node ID maps directly to an IPv6 link-local suffix at edge gateways, enabling IP-layer integration without transmitting 128-bit addresses over the constrained radio link.

The 8-bit address space supports up to 254 unicast nodes per RF cluster, which is sufficient for the dense but bounded deployments targeted by this work (e.g., agricultural sensor arrays, UAV swarm sub-clusters). For applications requiring broader network reach, Section VII foresees integration of full IPv6 routing, where the 8-bit node identifier serves as a local cluster suffix rather than a global unique identifier. Note that the 254 addressable nodes operate independently of the receiver's concurrency limit ($\text{MAX_CONCURRENT_SESSIONS} = 10$), which governs the maximum number of simultaneously active reassembly sessions.

The exact byte alignment and offsets of the 9-byte header packet structure are presented in Table III and Figure 1.

TABLE III
LAMP OTA PACKET LAYOUT AND BYTE OFFSETS

Offset	Field Name	Type	Description
0	srcAddr	uint8_t	Source Node ID (0x01–0xFE)
1	dstAddr	uint8_t	Destination Node ID (0x01–0xFE, 0xFF)
2–3	messageId	uint16_t	Message Sequence ID
4	totalChunks	uint8_t	Total Count ($1 \leq N \leq 255$; $N = 0$ invalid)
5	chunkIndex	uint8_t	Index ($0 \leq i < N$)
6	payloadSize	uint8_t	Chunk Payload ($0 \leq L \leq 244$ B; $L = 0$ for control)
7	flags	uint8_t	Bitfield (ACK_REQ, CONN_REQ)
8	protocolVersion	uint8_t	Protocol Version (0x02)
9–L+8	payload[]	uint8_t[]	Application Data (L bytes)
L+9–L+10	crc	uint16_t	CRC-16 Checksum

The fields in the Logical Header are defined as follows:

- **srcAddr (1 byte)**: Source node identifier (0x01–0xFE, or 0x00 for unassigned).
- **dstAddr (1 byte)**: Destination node identifier (0x01–0xFE, or 0xFF for broadcast).

- **messageId (2 bytes)**: Unique message sequence identifier. All fragments composing the same segmented message share this ID. Active and completed sessions are uniquely keyed by the 4-tuple ($\text{srcAddr}, \text{dstAddr}, \text{messageId}, \text{version}$).
- **totalChunks (1 byte)**: The total count of chunks into which the original large payload is divided ($1 \leq \text{totalChunks} \leq 255$).
- **chunkIndex (1 byte)**: The 0-based index of this specific fragment ($0 \leq \text{chunkIndex} < \text{totalChunks}$).
- **payloadSize (1 byte)**: Payload data byte count ($0 \leq L \leq 244$ bytes; $L = 0$ for pure control frames). On standard SX1262 transceivers (256-byte FIFO, 255-byte maximum PHY frame), subtracting the 9-byte header and 2-byte CRC-16 yields $S_{\max} = 244$ bytes per chunk, resulting in a single-session capacity of $255 \times 244 = 62,220$ bytes (≈ 62.2 KB).
- **flags (1 byte)**: A control bitfield used to manage transmission and connection states:
 - **FLAG_ACK_REQ (0x04)**: Selective Acknowledgment requested from receiver.
 - **FLAG_ACK (0x08)**: SACK feedback packet containing binary bitmap of received chunks.
 - **FLAG_CONN_REQ (0x10)**: Connection Request control frame.
 - **FLAG_CONN_ACK (0x20)**: Connection Acknowledged control frame.
 - **FLAG_CONN_NACK (0x40)**: Connection Refused control frame.
- **protocolVersion (1 byte)**: Identifies protocol version (currently 2) to ensure backward compatibility.

In LAMPv2, broadcast addressing (0xFF) is strictly restricted to unreliable Best-Effort Datagrams (Mode 1). To avoid acknowledgment implosion and channel collapse, control frames (FLAG_ACK_REQ, SYN, SYN-ACK, CONN_NACK) are prohibited with $\text{dstAddr} = 0xFF$.

After the payload, a 16-bit CRC-16 (Modbus CRC-16, polynomial 0xA001, initialized to 0xFFFF) is appended. The checksum is computed over the concatenation of the 9-byte Logical Header and exactly payloadSize application-layer bytes; any trailing buffer bytes beyond payloadSize are excluded, so that partial final chunks are validated identically to full-size chunks and inter-implementation interoperability is

Byte 0	Byte 1	Byte 2 – 3	Byte 4	Byte 5	Byte 6	Byte 7	Byte 8
srcAddr	dstAddr	messageId	totalChunks	chunkIndex	payloadSize	flags	protocolVersion

Fig. 1. Byte layout of the compact 9-byte Logical Header incorporating Node Addressing.

preserved.

Table IV summarises all defined flag combinations and the behaviour of the packet parser upon receipt of an undefined value.

TABLE IV
LAMP FLAG FIELD: VALID COMBINATIONS AND PARSER BEHAVIOUR

Value	Semantic	Notes
0x00	Data chunk	Unreliable/reliable chunk, no ACK
0x04	Final chunk (ACK_REQ)	Receiver replies with SACK
0x08	SACK feedback	Payload = received-chunk bitmap
0x10	SYN	3WHS Step 1 (SynMetadata)
0x30	SYN-ACK	3WHS Step 2 (SynAckMetadata)
0x20	ACK	3WHS Step 3, no payload ($L = 0$)
0x40	CONN-NACK	Payload = rejection reason code
other	Undefined	Silently discarded

C. Selective Acknowledgment and Retransmission

In reliable mode, the protocol implements a SACK mechanism [16] to recover lost fragments with minimal additional airtime.

1) *Retransmission Protocol Interaction:* When a transmission is executed in reliable mode:

- 1) The transmitter segments the payload, serializes the chunks with `srcAddr` and `dstAddr`, and streams them sequentially. The final chunk index ($Index = totalChunks - 1$) is marked with the `FLAG_ACK_REQ` flag.
- 2) Upon receipt of the chunk requesting acknowledgment, the receiver evaluates its reassembly session. It identifies which fragments are still missing by checking its internal bitmask of received chunks.
- 3) The receiver responds by transmitting a SACK packet (`FLAG_ACK`) addressed back to the transmitter's `srcAddr`. The SACK carries a payload of $\lceil N/8 \rceil$ bytes containing a binary bitmap: chunk index i ($0 \leq i < N$) maps to bit position ($i \bmod 8$) of byte $\lfloor i/8 \rfloor$ (LSB-first ordering), where a 1-bit indicates successful reception and a 0-bit indicates a missing fragment.
- 4) The transmitter parses the received SACK bitmap. If some chunks are missing, it selectively retransmits only those lost chunks. The last retransmitted chunk has the `FLAG_ACK_REQ` flag set again to query the receiver.
- 5) This selective retry loop continues until the full message is reconstructed at the receiver, or the transmitter reaches the maximum retry count.

2) *Half-Duplex Hardware Switching Guard:* Because LoRa transceivers are half-duplex, a hardware transition period is required for the transmitter to switch from transmit (TX) mode to receive (RX) mode after sending a packet. While physical SX1262 transceiver state transitions take $< 100 \mu s$, the $\delta = 25$ ms guard delay accounts for physical transceiver

switching alongside software-level Real-Time Operating System (RTOS) task scheduling jitter, Serial Peripheral Interface (SPI) bus transaction queuing, and microcontroller wake-up latencies on embedded targets.

It is important to emphasize that all protocol timing parameters, guard delays (e.g., `SACK_PREAMBLE_GUARD_DELAY_MS`), and retry limits are fully decoupled into a centralized configuration namespace (`LoRaMultiPacketConfig`) and can be customized by integrators for any MCU/radio target. The default parameters presented in this study were specifically tuned around the Heltec Wi-Fi LoRa 32 V3 (ESP32-S3) hardware platform.

3) *Timeout and Fallback Recovery:* If the transmitter does not receive a SACK packet within the dynamic timeout window, it assumes the query or response was lost and retransmits the last transmitted chunk with the `FLAG_ACK_REQ` flag. The transmitter permits up to $R_{\max} = 5$ consecutive timeout events before aborting the session.

The timeout is computed adaptively to account for forward data frame airtime, receiver turnaround, and reverse SACK airtime:

$$T_{\text{timeout}} = \max\left(T_{\min}, \alpha \cdot \left(\hat{T}_{\text{OA},\text{DATA}} + \hat{T}_{\text{OA},\text{SACK}}\right) + \delta\right) \quad (3)$$

where $T_{\min} = 1000$ ms is the minimum guard, $\alpha = 1.5$ is a safety factor, $\hat{T}_{\text{OA},\text{DATA}}$ and $\hat{T}_{\text{OA},\text{SACK}}$ are the Times-on-Air for the data frame and reverse SACK frame queried from the radio driver [17], and $\delta = 25$ ms is the half-duplex switching and scheduling guard delay. Table V reports representative timeout values across spreading factors.

4) *Idealized Analytical Burst Reliability Baseline:* Under an independent frame loss probability p and assuming error-free control feedback, the probability that an individual chunk fails after $R_{\max}$ repair rounds is $p^{R_{\max}+1}$. For an N -chunk payload under this idealized i.i.d. baseline, the overall burst delivery success probability P_{success} is modeled as:

$$P_{\text{success}} = (1 - p^{R_{\max}+1})^N, \quad P_{\text{fail}} = 1 - P_{\text{success}} \quad (4)$$

For a lossy channel with a 20% frame error rate ($p = 0.20$) and $R_{\max} = 5$ (matching `LoRaMultiPacketConfig::MAX_RETRIES`), $P_{\text{success}} = (1 - 0.20^6)^{20} \approx 99.87\%$ for a 20-chunk transfer. In physical wireless environments with correlated shadowing and reverse-link fading, overall completion rates are evaluated empirically via the stochastic ns-3 campaign in Section VI.

D. Memory Bounding and RAM Footprint Analysis

To guarantee memory safety on constrained microcontrollers, the dynamic memory allocated for an active reassembly session at the receiver is strictly bounded by:

$$M_{\text{session}} = S_{\text{struct}} + N_{\text{chunks}} \cdot S_{\text{max}} + \left\lceil \frac{N_{\text{chunks}}}{8} \right\rceil \quad (5)$$

TABLE V
ADAPTIVE SACK TIMEOUT VALUES (BW 500 kHz, CR 4/5, 255-BYTE
MAX FRAME)

SF	$\hat{T}_{\text{DATA}}$ (ms)	$\alpha \sum \hat{T}_{\text{oA}} + \delta$ (ms)	T_{timeout} (ms)
7	99	188	1000
8	174	312	1000
9	313	542	1000
10	564	958	1000
11	1025	1720	1720
12	1886	3140	3140

where $S_{\text{struct}} \approx 48$ bytes represents the session tracking structure metadata, $S_{\text{max}} = 244$ bytes is the maximum chunk stride, and the final term represents the binary SACK bitmask.

For a representative 20-chunk payload (4.88 KB data size), a complete reassembly session requires ≈ 4.93 KB of RAM (totaling ≈ 49.3 KB across 10 concurrent 20-chunk sessions). On entry-level microcontrollers with limited SRAM (≤ 64 KB), the concurrency pool ceiling (`MAX_CONCURRENT_SESSIONS`) is tuned to 1–2 sessions in `LoRaMultiPacketConfig`, while runtime heap protection (`AvailMem()`) prevents allocation failures by rejecting new requests with `CONN_NACK(BUSY)` when memory is depleted.

E. Formal Protocol Algorithms

To ensure complete specification clarity and system-level reproducibility across different hardware platforms, Algorithm 1 and Algorithm 2 formally detail the transmitter-side segmentation pipeline and the receiver-side out-of-order reassembly engine, respectively.

F. Session Control and Handshake Protocol

To support stateful stream negotiation and protect resource-constrained receivers from buffer overflows, LAMPv2 introduces a 3WHS mechanism.

1) *Three-Way Handshake Procedure*: As illustrated in Figure 3(a), the connection lifecycle follows a three-step exchange:

- 1) **SYN**: The initiator sends a SYN control frame (`FLAG_CONN_REQ`) addressed to the target node's `dstAddr`, containing a 4-byte `SynMetadata` payload (`requestedPayloadSize`, `windowSize`, `timeoutMs`).
- 2) **SYN-ACK**: Upon receiving the SYN, the peer evaluates available memory and protocol version. If accepted, it responds with a SYN-ACK control frame (`FLAG_CONN_REQ | FLAG_CONN_ACK`) addressed back to the initiator. The receiver clamps `acceptedPayloadSize` to its maximum buffer capability.
- 3) **ACK**: The initiator parses the SYN-ACK, adopts the negotiated payload size, and transmits a final ACK frame (`FLAG_CONN_ACK`). Both nodes transition their internal state machine to `ESTABLISHED`.

G. Concurrency and Fault Handling

To ensure protocol stability in multi-node and lossy environments, LAMP incorporates the following fault-handling mechanisms:

- *Simultaneous Connection Resolution*: If two nodes transmit SYN frames to each other simultaneously, a deterministic tie-breaking rule resolves the conflict: the node with the lower numerical `NodeAddress` yields and transitions back to `CLOSED`, while the node with the higher `NodeAddress` proceeds with session setup.
- *Receiver Busy Rejection*: If a receiver receives a SYN frame while its active reassembly session is already locked by another node, it responds with a `CONN_NACK` frame (`FLAG_CONN_NACK`) specifying a `BUSY` status code, causing the requester to back off.
- *Source Address Filtering*: Initiators in `waitForSynAck` and `waitForSack` validate incoming control frames against the expected `srcAddr`, ignoring packets from unrelated third-party nodes.
- *Incomplete Handshake Timeout*: If a SYN-ACK or ACK of a 3WHS is lost in transit, the receiver in `SYN_RCVD` or initiator in `SYN_SENT` automatically aborts and resets to `CLOSED` after an inactivity timeout (`timeoutMs`).

IV. PRELIMINARY PROOF OF CONCEPT

An initial prototype using commercial Ebyte E220-series modules validated the core segmentation concept. Field trials confirmed feasibility but exposed three limitations: (1) the modules exposed a proprietary network layer that prevented direct control of physical-layer parameters; (2) a large, unoptimized application header inflated per-chunk airtime; and (3) no acknowledgment or retransmission mechanism was available.

These findings motivated the redesign on Semtech SX126x-class transceivers via the `RadioLib` driver, validated on Heltec Wi-Fi LoRa 32 V3 boards. The revised implementation described in the following sections, is fully hardware-agnostic through the HAL interface and incorporates the 9-byte Logical Header, CRC-16 integrity checks, and the SACK-based retransmission mechanism.

V. HARDWARE LEVEL IMPLEMENTATION

To validate the proposed protocol, a complete implementation of the LAMP library was developed in C++17. The design is modular, separating the transport state machine from physical hardware dependencies and parameter policies.

A. Centralized Configuration and Parameter Decoupling

To eliminate hardcoded constants and enable deterministic tuning across heterogeneous embedded targets, all protocol operational parameters are centralized within a dedicated compile-time configuration namespace (`LoRaMultiPacketConfig`). This design encapsulates three primary domains:

- **Memory and Concurrency Limits**: Governs maximum concurrent reassembly sessions (`MAX_CONCURRENT_SESSIONS = 10`, reducible to

Algorithm 1 LAMP Payload Transmission and SACK-ARQ (Transmitter)

Require: Payload P (size L), $dstAddr$, Max Chunk Size S , Mode M
Ensure: Delivery Status (OK or FAIL)

```

1:  $N \leftarrow \lceil L/S \rceil$ 
2: if  $N > 255$  then
3:   return ERR_PAYLOAD_TOO_LARGE
4: end if
5:  $msgId \leftarrow \text{GetNextMsgId}()$ 
6: for  $i \leftarrow 0$  to  $N - 1$  do
7:    $cLen \leftarrow (i == N - 1) ? (L - i \cdot S) : S$ 
8:    $flags \leftarrow (M \in \{\text{Mode 3}, \text{Mode 4}\} \text{ and } i == N - 1) ? \text{FLAG\_ACK\_REQ} : 0x00$ 
9:    $hdr \leftarrow \text{BuildHdr}(src, dst, msgId, N, i, cLen, flags)$ 
10:   $crc \leftarrow \text{CalcCRC16}(hdr, P[i \cdot S : i \cdot S + cLen])$ 
11:   $F_i \leftarrow \text{SerializeFrame}(hdr, P[i \cdot S : i \cdot S + cLen], crc)$ 
12:   $\text{Transmit}(F_i)$ 
13: end for
14: if  $M \in \{\text{Mode 1}, \text{Mode 2}\}$  then
15:   return OK ▷ Unreliable delivery complete
16: end if
17:  $retries \leftarrow 0$ ;  $lastIdx \leftarrow N - 1$ 
18: while  $retries \leq R_{\max}$  do
19:    $sack \leftarrow \text{WaitForSack}(T_{\text{timeout}})$ 
20:   if  $sack == \text{TIMEOUT}$  then
21:     if  $retries == R_{\max}$  then
22:       break
23:     end if
24:      $\text{Transmit}(F_{lastIdx})$ ;  $retries \leftarrow retries + 1$ 
25:   else if  $\text{IsAllAacked}(sack.bitmap)$  then
26:     return OK ▷ Full payload delivery confirmed
27:   else
28:     if  $retries == R_{\max}$  then
29:       break
30:     end if
31:      $missing \leftarrow \text{MissingChunks}(sack.bitmap)$ 
32:     for  $idx \in missing$  do
33:        $cLen \leftarrow (idx == N - 1) ? (L - idx \cdot S) : S$ 
34:        $flags \leftarrow (idx == \text{Last}(missing)) ? \text{FLAG\_ACK\_REQ} : 0x00$ 
35:        $hdr \leftarrow \text{BuildHdr}(src, dst, msgId, N, idx, cLen, flags)$ 
36:        $crc \leftarrow \text{CalcCRC16}(hdr, P[idx \cdot S : idx \cdot S + cLen])$ 
37:        $F_{idx} \leftarrow \text{SerializeFrame}(hdr, P[idx \cdot S : idx \cdot S + cLen], crc)$ 
38:        $\text{Transmit}(F_{idx})$ 
39:     end for
40:      $lastIdx \leftarrow \text{Last}(missing)$ ;  $retries \leftarrow retries + 1$ 
41:   end if
42: end while
43: return ERR_MAX_RETRIES_EXCEEDED

```

1–2 on memory-constrained Cortex-M0/AVR devices), physical FIFO buffer bounds ($\text{PHY_BUFFER_SIZE} = 256$ bytes, reflecting transceiver hardware FIFO capacity supporting the 255-byte maximum LoRa PHY payload cap), and completed sequence deduplication depth ($\text{MAX_COMPLETED_HISTORY} = 16$).

- **Timing and Retransmission Policies:** Configures the default retry ceiling ($R_{\max} = 5$), minimum SACK timeout floor ($T_{\min} = 1000$ ms, tunable for low-latency P2P links), safety multiplier ($\alpha = 1.5$), and half-duplex switching guard ($\delta = 25$ ms).
- **Session Lifecycle Thresholds:** Defines handshake initiation timeouts ($\text{DEFAULT_SYN_TIMEOUT_MS} = 3000$ ms), connection inactivity teardown periods ($\text{DEFAULT_CONN_INACTIVITY_TIMEOUT_MS} = 15000$ ms), and stale session pruning timers

Algorithm 2 Out-of-Order Reassembly and SACK Generation (Receiver)

Require: Received Frame F , Fixed Chunk Stride $S_{\max}$
Ensure: Reassembled Payload P or Drop Status

```

1:  $(hdr, payload, crc) \leftarrow \text{ParseFrame}(F)$ 
2: if  $\text{ValidateCRC}(hdr, payload, crc) == \text{FAIL}$  then
3:    $\text{Drop}(F)$ ; return
4: end if
5: if  $hdr.dstAddr \neq myAddr$  and  $hdr.dstAddr \neq 0xFF$  then
6:    $\text{Drop}(F)$ ; return
7: end if
8: if  $hdr.dstAddr == 0xFF$  and  $(hdr.flags \& \text{FLAG\_ACK\_REQ})$  then
9:    $\text{Drop}(F)$ ; return
10: end if
11: if  $hdr.totalChunks == 0$  or  $hdr.chunkIndex \geq hdr.totalChunks$  then
12:    $\text{Drop}(F)$ ; return
13: end if
14: if  $\text{IsCompleted}(hdr.srcAddr, hdr.msgId)$  then
15:   if  $hdr.flags \& \text{FLAG\_ACK\_REQ}$  then
16:      $\text{SendTerminalSack}(hdr.srcAddr, hdr.msgId)$ 
17:   end if
18:    $\text{Drop}(F)$ ; return
19: end if
20:  $session \leftarrow \text{FindSession}(hdr.srcAddr, hdr.msgId)$ 
21: if  $session == \text{NULL}$  then
22:   if  $\text{AvailMem}() < hdr.totalChunks \cdot S_{\max}$  then ▷ Safe stride check
23:      $\text{SendNack}(hdr.srcAddr, \text{BUSY})$ ; return
24:   end if
25:    $session \leftarrow \text{AllocBuf}(hdr.totalChunks, S_{\max})$ 
26: end if
27:  $idx \leftarrow hdr.chunkIndex$ 
28: if  $\text{IsBitSet}(session.bitmap, idx)$  then
29:    $\text{Drop}(F)$  ▷ Duplicate chunk, ignore write
30: else
31:    $\text{WriteChunk}(session.buffer, idx \cdot S_{\max}, payload)$ 
32:    $\text{SetBit}(session.bitmap, idx)$ ;  $session.count \leftarrow session.count + 1$ 
33: end if
34: if  $hdr.flags \& \text{FLAG\_ACK\_REQ}$  then
35:    $\text{SendSack}(hdr.srcAddr, hdr.msgId, session.bitmap)$ 
36: end if
37: if  $session.count == hdr.totalChunks$  then
38:    $P \leftarrow \text{ExtractPayload}(session.buffer)$ 
39:    $\text{DeliverToApp}(P)$ ;  $\text{MarkCompleted}(hdr.srcAddr, hdr.msgId)$ ;  $\text{FreeSession}(session)$ 
40: end if

```

($\text{PRUNE_TIMEOUT_MS} = 15000$ ms).

Complementing the compile-time configuration, the 3WHS mechanism enables dynamic run-time parameter negotiation (SynMetadata), providing a two-tier flexibility model that adapts to variable channel conditions and peer memory capabilities.

B. Hardware Abstraction Layer (HAL)

To keep the protocol core hardware-agnostic, a clean HAL interface (RadioLibHal) was implemented. On the physical nodes, the concrete class EspHal wraps ESP32-S3 hardware resources: Serial Peripheral Interface (SPI) Communication (register configuration and FIFO buffer transfers for the Semtech SX1262), General-Purpose Input/Output (GPIO) Interrupts (NSS, RST, BUSY, and the DIO1 RX/TX comple-




Fig. 2. Finite State Machine (FSM) specifications for connection session lifecycle and data transfer modes.




Fig. 3. Sequence interaction flows for 3-Way Handshake (3WHS) connection setup with dynamic metadata negotiation and rejection handling.




Fig. 4. End-to-end data streaming flow, frame loss detection, and SACK-driven selective retransmission interaction.

tion interrupt pin), and Timing Primitives (`millis()` and `delay()`).

C. Host-Based Unit Testing

Following a TDD methodology, the protocol core was validated in isolation on a desktop host using the **Unity** framework in PlatformIO, verifying: (i) Modbus CRC-16 calculation on headers and partial/padded payloads, (ii) reassembler state transitions under out-of-order and duplicate chunk arrivals, and (iii) deterministic pruning of stale session descriptors.

D. Physical Benchmarking Setup

Physical validation was conducted using two Heltec Wi-Fi LoRa 32 V3 development boards (ESP32-S3 + Semtech SX1262) configured with: Carrier Frequency 869.525 MHz, SF7, BW 500.0 kHz, CR 4/5, TX Power +10 dBm, and 8-symbol preamble.

1) *Regulatory Compliance and Regional Deployment*: It is important to clarify the relationship between the LAMP transport protocol and RF regulatory constraints. LAMP operates strictly at the *transport layer* and is agnostic to regional carrier frequency, channel width, and RF power limits; compliance with these constraints is the responsibility of the underlying MAC/PHY layer.

Under European ETSI EN 300 220-2 regulations [6], the 869.400–869.650 MHz sub-band (Band P) allocates a 250 kHz total frequency block permitting up to 10% duty cycle and an output power up to 500 mW ERP (27 dBm). For field deployments strictly adhering to this sub-band, channel bandwidths of ≤ 250 kHz (specifically 125 kHz or 250 kHz raster centered at 869.525 MHz) are selected. To characterize the protocol's upper theoretical throughput limits under maximum channel stress, our simulation campaign and laboratory benchmarks evaluated a wideband configuration ($BW = 500.0$ kHz, matching standard US915 / AU915 wideband LoRa allocations) on an experimental carrier.

Additionally, the reference C++ library provides opt-in MAC-layer helpers, including SX1262-native Channel Activity Detection (CAD) with randomized exponential backoff (CSMA-CA) and duty-cycle pacing to mitigate channel contention. These features are decoupled from the transport core and offered as convenience primitives for integrators targeting shared ISM spectrum.

VI. PERFORMANCE EVALUATION AND EXPERIMENTAL METHODOLOGY

To evaluate the reliability, throughput, and communication limits of LAMP, we adopt a two-stage evaluation methodology: a large-scale stochastic network simulation campaign executed in ns-3, complemented by functional hardware testbed validation on microcontroller nodes and a planned outdoor field benchmarking campaign.

A. Network Simulation Benchmarking (ns-3)

To characterize protocol performance under controlled fading channels across extensive distance ranges, we developed a full C++ simulator integrated into the ns-3 framework (ns3-lorawan). The simulator reproduces the exact LAMPv2 frame structure (9-byte Logical Header + 2-byte CRC-16), all four operational modes, the 3WHS handshake, and selective bit-vector SACK retransmissions over a Nakagami- m fast-fading propagation loss model calibrated for the Semtech SX1262 transceiver (+24.2 dBm EIRP, 869.525 MHz), with shape factors: $m = 3.0$ for LoS (strong specular component), $m = 1.5$ for Urban (moderate multipath), and $m = 1.0$ for Rural (Rayleigh fading).

1) *Simulation Campaign Hierarchy*: To achieve statistical convergence, an extensive Monte Carlo campaign comprising 117,050 total independent simulation runs was conducted. The experimental space spans 3 RF propagation environments (LoS, Urban, Rural), 3 Spreading Factors (SF7, SF9, SF10), 4 operational modes (Mode 1 to Mode 4), and 50 independent random seeds per parameter configuration across an adaptive distance grid extending dynamically up to 23.5 km with

early-stopping thresholds once reception falls to 0% across consecutive steps. Each curve point plotted in Figs. 5 and 6 represents the sample mean over 50 independent Monte Carlo seed replications for a representative $N = 20$ -chunk payload (4.88 KB data size), with shaded ribbons denoting 95% confidence intervals.

2) *Analytical Reliability Bound under Stated Assumptions*: To model payload completion probabilities analytically, we define a theoretical baseline under stated hypotheses:

- 1) *Independent Loss Model*: Channel frame loss is modeled as an independent and identically distributed (i.i.d.) Bernoulli error process with packet error rate $p \in [0, 1)$ and error-free control feedback ($p_{\text{SACK}} = 0$). While the i.i.d. Bernoulli model provides a tractable analytical baseline, physical LPWAN channels typically exhibit temporally correlated fading.
- 2) *Bounded Repair Rounds*: Selective SACK retransmissions are bounded by the protocol's maximum repair round budget $R_{\text{max}} = 5$ (matching `LoRaMultiPacketConfig::MAX_RETRIES`).

Under these idealized assumptions, the analytical upper bound for successfully reassembling an N -fragment payload across R_{max} retries matches Eq. (4):

$$P_{\text{complete}}(N, p, R_{\text{max}}) = (1 - p^{R_{\text{max}}+1})^N \quad (6)$$

Because Algorithm 1 enforces a global session repair round budget ($R_{\text{max}} = 5$) across all missing fragments rather than independent per-chunk retries, Eq. (6) serves as an optimistic theoretical baseline. Empirical multi-chunk delivery dynamics under correlated channel fading and joint forward-reverse link losses are directly evaluated via the stochastic ns-3 campaign. It is crucial to clarify that when reliable modes (Mode 3 and Mode 4) sustain a 100% Application Payload Delivery Ratio (PDR_{app}) in Fig. 5, this does not indicate physical channel reliability or zero-loss propagation. Rather, it demonstrates that the selective bit-vector SACK-ARQ mechanism successfully recovers lost fragments within the $R_{\text{max}} = 5$ retry budget prior to session timeout.

In LoS conditions, SF10 sustains complete payload reassembly up to 19.1 km, exhibiting a steep drop-off between 20.0 km and 23.1 km as frame loss probability exceeds retransmission recovery budgets. Conversely, best-effort unreliable modes (Mode 1 and Mode 2) experience progressive packet drop rates as spatial fading degrades signal-to-noise ratios.

Fig. 6 depicts effective application-layer Goodput (bps). While best-effort streaming (Mode 2) achieves peak throughput in pristine channels by omitting acknowledgment latency, Mode 4 (*Reliable Stream*) preserves sustained throughput near physical channel capacity across the operational range, degrading only near receiver sensitivity thresholds.

B. Planned Real-World Field Benchmarking Campaign

While hardware testbed experiments on two Heltec Wi-Fi LoRa 32 V3 nodes have validated the functional correctness, C++ memory safety, 3WHS state machine transitions, and out-of-order reassembly of LAMP, a systematic outdoor field campaign is planned as future experimental work to




Fig. 5. Packet Delivery Ratio (PDR_{app} %) for a fixed $N = 20$ -chunk payload transfer (4.88 KB application data) as a function of distance across Spreading Factors (SF7, SF9, SF10) in Line-of-Sight (LoS) propagation. Solid lines denote mean delivery rates across 50 independent Monte Carlo random seeds; shaded ribbons represent 95% confidence intervals.




Fig. 6. Effective application-layer throughput (Goodput in bps) for an $N = 20$ -chunk payload transfer across Spreading Factors and LAMP operational modes in LoS. Shaded ribbons denote 95% confidence intervals across 50 Monte Carlo seeds.

record real-world empirical performance under environmental interference, multipath reflections, and temperature-induced frequency drift.

1) Planned Experimental Matrix and Evaluation Metrics:

The planned outdoor physical validation strategy follows established LPWAN characterization methodologies [10]. It will sweep transmission distances (1 to 5000 meters across LoS and Non-Line-of-Sight (NLoS) conditions) and payload sequence sizes (1 to 100 chunks, ≈ 24.6 KB) under standard radio configurations (SF7, BW 500 kHz, CR 4/5, 10 dBm TX power).

During the planned outdoor campaign, four core performance metrics will be recorded across 5 independent trials per operational mode: (i) **Application Payload Delivery Ratio (PDR_{app})**, measuring successful reassembly rate; (ii) **End-to-End Latency**, from initial serialization to reassembly completion; (iii) **Effective Throughput (Goodput)**, calculated as delivered payload bits over elapsed duration; and (iv) **Retransmission Overhead**, tracking extra frames and SACK packets relative to unreliable baselines.

VII. CONCLUSION AND ARCHITECTURAL ROADMAP

This paper presents the design, implementation, network simulation, and hardware validation of LAMP, a lightweight and reliable transport-layer architecture optimized for transmitting segmented large payloads over point-to-point and local multi-node LoRa networks. By replacing traditional Stop-and-Wait per-packet acknowledgments with selective retransmissions requested via cumulative binary bitmaps, and eliminating redundant SOM/EOM markers in favor of implicit index derivation, the protocol achieves high reliability under lossy conditions with minimal airtime overhead and duty-cycle footprint.

A. Protocol Scope and Limitations

To provide a transparent systems assessment, three inherent architectural boundaries of baseline LAMPv2 are identified:

- 1) *Session Capacity Boundary*: A single transport session is bounded to $255 \text{ chunks} \times 244 \text{ B} = 62,220 \text{ bytes}$ ($\approx 62.2 \text{ KB}$). Larger files require application-layer multi-session chunking.
- 2) *Pessimistic Buffer Pre-allocation*: To guarantee memory safety against out-of-order partial tail chunks arriving first, the receiver allocates contiguous memory based on $N \times S_{\max}$ upfront. On RAM-constrained micro-controllers ($\leq 32 \text{ KB SRAM}$), the concurrency pool (`MAX_CONCURRENT_SESSIONS`) must be set to 1–2 sessions.
- 3) *Unauthenticated Baseline 3WHS*: Baseline LAMPv2 incorporates CRC-16 for physical frame integrity but omits cryptographic nonces. Consequently, adversarial replays of stale SYN frames could exhaust active session descriptors. This is mitigated via the opt-in Level-2 Security Profile (AES-128-GCM / HMAC) detailed below.

B. Architectural Roadmap

To guide the long-term progression of LAMP from point-to-point data links to decentralized multiagent mesh networks, we define a three-level architectural roadmap:

- **Level 1 – Addressed Transport & Polite Spectrum Access (Current Stack)**: 8-bit addressing (`srcAddr`, `dstAddr`), 9-byte header, four operational modes, dynamic 3WHS, and opt-in CAD/CSMA-CA under ETSI EN 300 220-1 [6].
- **Level 2 – Adaptive Spectrum, Scheduling & Security (Near-Term)**: Dynamic ADR tuning via RSSI/SNR metrics in SACK feedback, time-slotted channel access (TDM) [15], and an optional transport security profile (AES-128-GCM / HMAC-SHA256 with 32-bit replay counters).
- **Level 3 – Multi-Hop Mesh & IPv6 Extension (Long-Term Vision)**: Dynamic multi-hop routing (AODV-Light / RPL adaptation) with hop-by-hop vs. end-to-end SACK

reliability [13], [14], and 8-bit node ID mapping for IPv6 link-local bridging (fe80::/64).

DECLARATIONS

Data and Code Availability: The complete C++17 library source code, HAL drivers, Unity test suites, and ns-3 simulation scripts are publicly available on GitHub at <https://github.com/aleemont/LoRaMultiPacket> under the MIT license.

Ethical Statement & AI Tool Usage: This study did not involve human or animal subjects. AI assistant tools were utilized strictly for code formatting, micro-typography auditing, and manuscript style calibration in compliance with IEEE publication ethics guidelines.

Conflict of Interest: The authors declare that they have no competing financial or personal interests that could influence the work reported in this paper.

REFERENCES

- [1] M. Radeta, J. Pestana, P. Abreu, R. Freitas, F. Silva, D. Vieira, R. Prieto, M. Fernandez, F. Alves, T. Dellinger, S. Neves, and E. Delory, "Triton—open telemetry and location estimation for marine monitoring based on iot and lora," *IEEE Journal of Oceanic Engineering*, vol. 50, no. 2, pp. 1244–1258, 2025.
- [2] T. Elshabrawy and J. Al-Ali, "Lora technology demystified: From link behavior to cell capacity," *IEEE Transactions on Wireless Communications*, vol. 17, no. 10, pp. 6611–6621, 2018.
- [3] F. Adelantado, X. Vilajosana, P. Tuset-Peiro, B. Martinez, J. Melia-Segui, and T. Watteyne, "Understanding the limits of lorawan implementations," *IEEE Communications Magazine*, vol. 55, no. 9, pp. 34–40, 2017.
- [4] S. De, H. M. Jalajamony, S. Adhinarayanan, S. Joshi, H. Upadhyay, and R. Fernandez, "Multimedia transmission over lora networks for iot applications: A survey of strategies, deployments, and open challenges," *Sensors*, vol. 25, no. 23, p. 7128, 2025.
- [5] P. J. Marcellis, V. S. Rao, and R. V. Prasad, "Dare: Data recovery through application layer coding for lorawan," in *Proceedings of the 2017 IEEE/ACM Second International Conference on Internet-of-Things Design and Implementation (IoTDI)*, 2017.
- [6] *Short Range Devices (SRD) operating in the frequency range 25 MHz to 1 000 MHz; Part 1: Technical characteristics and methods of measurement*, European Telecommunications Standards Institute (ETSI) Std. ETSI EN 300 220-1 V3.1.1, 2017.
- [7] A. Minaburo, L. Toutain, C. Gomez, D. Barthel, and J. C. Zúñiga, "SCHC: Static context header compression and fragmentation for LP-WAN," Internet Engineering Task Force (IETF), RFC 8724, April 2020.
- [8] C. Gomez and J. Crowcroft, "Static context header compression (SCHC) over LoRaWAN," Internet Engineering Task Force (IETF), RFC 9011, April 2021.
- [9] LoRa Alliance Technical Committee, "LoRaWAN fragmented data block transport specification," LoRa Alliance, Technical Specification TS004-2.0.0, 2022.
- [10] M. Cattani, C. A. Boano, and K. Römer, "An experimental evaluation of the reliability of lora long-range low-power wireless communication," *Journal of Sensor and Actuator Networks*, vol. 6, no. 2, p. 7, 2017.
- [11] M. Bor, U. Roedig, T. Voigt, and J. M. Alonso, "Do lora low-power wide-area networks scale?" in *Proceedings of the 19th ACM International Conference on Modeling, Analysis and Simulation of Wireless and Mobile Systems (MSWiM)*, 2016, pp. 59–66.
- [12] T. Chen, D. Eager, and D. Makaroff, "Efficient image transmission using lora technology in agricultural monitoring iot systems," in *2019 IEEE International Conference on Internet of Things (iThings) and IEEE Green Computing and Communications (GreenCom) and IEEE Cyber, Physical and Social Computing (CPSCom) and IEEE Smart Data (SmartData)*, 2019, pp. 937–944.
- [13] J. Kim, J. Jenkins, J. Seol, and M. Kwak, "Extending data transmission in the multi-hop lora network," *Issues in Information Systems*, vol. 23, no. 3, pp. 292–304, 2022.
- [14] A. W.-L. Wong, S. L. Goh, M. K. Hasan, and S. Fattah, "Multi-hop and mesh for lora networks: Recent advancements, issues, and recommended applications," *ACM Computing Surveys*, vol. 56, no. 6, pp. 1–43, 2024.
- [15] C.-C. Wei, J.-K. Huang, C.-C. Chang, and K.-C. Chang, "The development of lora image transmission based on time division multiplexing," in *2020 International Symposium on Computer, Consumer and Control (IS3C)*, 2020, pp. 323–326.
- [16] A. Calabrò, E. Marchetti, and M. T. Paratore, "Evaluating use of arq strategies in communication protocols for search and rescue," in *Proceedings of the 21st International Conference on Web Information Systems and Technologies (WEBIST)*. SCITEPRESS, 2025, pp. 59–70.
- [17] S. Corporation, "Lora modem designer's guide," Semtech Corporation, Tech. Rep. Application Note AN1200.13, 2015.